

# Introduction to Coccinelle

Julia Lawall (Inria/LIP6)

<http://coccinelle.lip6.fr>

October 31, 2013

## Common programming problems

- Programmers don't really understand how C works.
  - `!e1 & e2` does a bit-and with 0 or 1.
- A simpler API function exists, but not everyone uses it.
  - Mixing different functions for the same purpose is confusing.
- A function may fail, but the call site doesn't check for that.
  - A rare error case will cause an unexpected crash.
- Etc.

Need for pervasive code changes.

## Example: Bad bit-and

```
if (!dma_cntrl & DMA_START_BIT) {  
    BCMLOG(BCMLOG_DBG, "Already Stopped\n");  
    return BC_STS_SUCCESS;  
}
```

From `drivers/staging/crystalhd/crystalhd_hw.c`

## Example: Inconsistent API usage

drivers/mtd/nand/r852.c:

```
if (!bounce) {
    dev->phys_dma_addr =
        pci_map_single(dev->pci_dev, (void *)buf, R852_DMA_LEN,
            (do_read ? PCI_DMA_FROMDEVICE : PCI_DMA_TODEVICE));

    if (pci_dma_mapping_error(dev->pci_dev, dev->phys_dma_addr))
        bounce = 1;
}
```

drivers/mtd/nand/denali.c:

```
denali->buf.dma_buf =
    dma_map_single(&dev->dev, denali->buf.buf, DENALI_BUF_SIZE,
        DMA_BIDIRECTIONAL);
if (dma_mapping_error(&dev->dev, denali->buf.dma_buf)) ...
pci_set_master(dev);
...
ret = pci_request_regions(dev, DENALI_NAND_NAME);
```

## Example: Missing error check

```
alloc = kmalloc(sizeof *alloc, GFP_KERNEL);  
INIT_LIST_HEAD(&intmem_allocations);  
intmem_virtual = ioremap(MEM_INTMEM_START + RESERVED_SIZE,  
                          MEM_INTMEM_SIZE - RESERVED_SIZE);  
  
initiated = 1;  
alloc->size = MEM_INTMEM_SIZE - RESERVED_SIZE;
```

From arch/cris/arch-v32/mm/intmem.c

## Our goals

- Automatically **find** code containing bugs or defects, or requiring collateral evolutions.
- Automatically **fix** bugs or defects, and perform collateral evolutions.
- Provide a system that is **accessible** to software developers.

# Requirements for automation

The ability to abstract over irrelevant information:

- `if (!dma_cntrl & DMA_START_BIT) { ... }:`  
    `dma_cntrl` is not important.

The ability to match scattered code fragments:

- `kmalloc` may be far from the first dereference.

The ability to transform code fragments:

- Replace `pci_map_single` by `dma_map_single`, or vice versa.

# Coccinelle

Program matching and transformation for unprocessed C code.

Fits with the existing habits of C programmers.

- C-like, patch-like notation

Semantic patch language (SmPL):

- **Metavariables** for abstracting over subterms.
- “...” for abstracting over code sequences.
- Patch-like notation ( $-/+$ ) for expressing transformations.



## The !& problem

**The problem:** Combining a boolean (0/1) with a constant using & is usually meaningless:

```
if (!dma_cntrl & DMA_START_BIT) {  
    BCMLOG(BCMLOG_DBG, "Already Stopped\n");  
    return BC_STS_SUCCESS;  
}
```

**The solution:** Add parentheses.

**Our goal:** Do so automatically for any expression **E** and constant **C**.

## A semantic patch for the !& problem

@@

expression E;

constant C;

@@

- !E & C

+ !(E & C)

Two parts per rule:

- Metavariable declaration
- Transformation specification

A semantic patch can contain multiple rules.

# Metavariable types

Surrounded by @@ @@.

- expression, statement, type, constant, local idexpression
- A type from the source program
- iterator, declarer, iterator name, declarer name, typedef

## Transformation specification

- - in the leftmost column for something to remove
- + in the leftmost column for something to add
- \* in the leftmost column for something of interest
  - Cannot be used with + and -.
- Spaces, newlines irrelevant.

## Exercise 1

1. Create a file ex1.cocci containing the following:

```
@@
```

```
expression E;
```

```
constant C;
```

```
@@
```

```
- !E & C
```

```
+ !(E & C)
```

2. Run spatch: `spatch --sp-file ex1.cocci --dir linux-3.2/drivers/staging/crystalhd`
3. Did your semantic patch do everything it should have?
4. Did it do something it should not have?

## Exercise 2

Some code contains a cast on the result of `kmalloc`. For example:

```
info->RegsBuf = (unsigned char *)  
    kmalloc(sizeof(info->ATAREgs), GFP_KERNEL);
```

If the destination of the returned value has pointer type, this cast is not needed.

1. Complete the following semantic patch to remove this unnecessary cast.

```
@@ expression * e; expression arg1, arg2; type T; @@
```

[fill it in]

2. Test your semantic patch on the code in `linux-3.2/drivers/isdn`
3. Are you satisfied with the appearance of the results? If not, try to improve it.

## Practical issues

To check that your semantic patch is valid:

```
spatch --parse-cocci mysp.cocci
```

To run your semantic patch:

```
spatch --sp-file mysp.cocci file.c  
spatch --sp-file mysp.cocci --dir directory
```

To understand why your semantic patch didn't work:

```
spatch --sp-file mysp.cocci file.c --debug
```

If you don't need to include header files:

```
spatch --sp-file mysp.cocci --dir directory  
--no-includes --include-headers
```

## More practical issues

Put the interesting output in a file:

```
spatch ... > output.patch
```

Omit the uninteresting output:

```
spatch --very-quiet ...
```

The source code:

```
/usr/src/linux-source-3.2/scripts/coccinelle/
```

These slides:

```
http://pagesperso-systeme.lip6.fr/Julia.Lawall/tutorial/part1.pdf
```



# Inconsistent API usage

Do we need this function?

```
static inline dma_addr_t
pci_map_single(struct pci_dev *hwdev, void *ptr, size_t size,
               int direction)
{
    return dma_map_single(hwdev == NULL ? NULL : &hwdev->dev, ptr,
                          size, (enum dma_data_direction)direction);
}
```

# The use of `pci_map_single`

The code:

```
dev->phys_dma_addr =  
    pci_map_single(dev->pci_dev, (void *)buf, R852_DMA_LEN,  
        (do_read ? PCI_DMA_FROMDEVICE : PCI_DMA_TODEVICE));
```

would be more uniform as:

```
dev->phys_dma_addr =  
    dma_map_single(&dev->pci_dev->dev, (void *)buf, R852_DMA_LEN,  
        (do_read ? DMA_FROM_DEVICE : DMA_TO_DEVICE));
```

Issues:

- Change function name.
- Add field access to the first argument.
- Rename the fourth argument.

## pci\_map\_single: Example and definitions

Commit b0eb57cb

```
- rbi->dma_addr = pci_map_single(adapter->pdev,  
+ rbi->dma_addr = dma_map_single(  
+     &adapter->pdev->dev,  
      rbi->skb->data, rbi->len,  
      PCI_DMA_FROMDEVICE);
```

### PCI constants

```
/* This defines the direction arg  
   to the DMA mapping routines. */  
#define PCI_DMA_BIDIRECTIONAL    0  
#define PCI_DMA_TODEVICE         1  
#define PCI_DMA_FROMDEVICE       2  
#define PCI_DMA_NONE             3
```

### DMA constants

```
enum dma_data_direction {  
    DMA_BIDIRECTIONAL = 0,  
    DMA_TO_DEVICE = 1,  
    DMA_FROM_DEVICE = 2,  
    DMA_NONE = 3,  
};
```

## pci\_map\_single: First attempt

Outline of a semantic patch, including the patch example:

@@

@@

```
- rbi->dma_addr = pci_map_single(adapter->pdev,  
+ rbi->dma_addr = dma_map_single(  
+     &adapter->pdev->dev,  
      rbi->skb->data, rbi->len,  
      PCI_DMA_FROMDEVICE);
```

## pci\_map\_single: First attempt

Eliminate irrelevant code:

@@

@@

```
- pci_map_single(adapter->pdev,  
+ dma_map_single(  
+     &adapter->pdev->dev,  
      rbi->skb->data, rbi->len,  
      PCI_DMA_FROMDEVICE)
```

## pci\_map\_single: First attempt

Abstract over subterms:

@@

expression E1,E2,E3;

@@

```
- pci_map_single(E1,  
+ dma_map_single(  
+   &E1->dev,  
      E2, E3,  
      PCI_DMA_FROMDEVICE)
```

## pci\_map\_single: First attempt

Rename the fourth argument:

@@

expression E1,E2,E3;

@@

```
- pci_map_single(E1,  
+ dma_map_single(  
+     &E1->dev,  
      E2, E3,  
-     PCI_DMA_FROMDEVICE)  
+     DMA_FROM_DEVICE)
```

## pci\_map\_single: Second attempt

Need to consider all direction constants.

```
@@ expression E1,E2,E3; @@  
- pci_map_single(E1,  
+ dma_map_single(&E1->dev,  
    E2, E3,  
-    PCI_DMA_FROMDEVICE)  
+    DMA_FROM_DEVICE)
```

```
@@ expression E1,E2,E3; @@  
- pci_map_single(E1,  
+ dma_map_single(&E1->dev,  
    E2, E3,  
-    PCI_DMA_TODEVICE)  
+    DMA_TO_DEVICE)
```

Etc. Four rules in all.



## pci\_map\_single: Third attempt

Avoid code duplication: Use a disjunction.

```
@@ expression E1,E2,E3; @@  
- pci_map_single(E1,  
+ dma_map_single(&E1->dev,  
    E2, E3,  
(  
-     PCI_DMA_BIDIRECTIONAL  
+     DMA_BIDIRECTIONAL  
|  
-     PCI_DMA_TODEVICE  
+     DMA_TO_DEVICE  
|  
-     PCI_DMA_FROMDEVICE  
+     DMA_FROM_DEVICE  
|  
-     PCI_DMA_NONE  
+     DMA_NONE_DEVICE  
)  
)
```

## pci\_map\_single: Fourth attempt

```
@@ expression E1,E2,E3,E4; @@
```

```
- pci_map_single(E1,  
+ dma_map_single(&E1->dev,  
    E2, E3, E4)
```

```
@@ expression E1,E2,E3; @@
```

```
dma_map_single(E1, E2, E3,  
(  
-     PCI_DMA_BIDIRECTIONAL  
+     DMA_BIDIRECTIONAL  
|  
-     PCI_DMA_TODEVICE  
+     DMA_TO_DEVICE  
|  
-     PCI_DMA_FROMDEVICE  
+     DMA_FROM_DEVICE  
|  
-     PCI_DMA_NONE  
+     DMA_NONE_DEVICE  
)  
)
```

## Exercise 3

1. Implement some version of the semantic patch for converting calls to `pci_map_single` to calls to `dma_map_single`.
2. Test your implementation on the directory `linux-3.2/drivers/net/ethernet`.
3. Implement both the third version and the fourth version. Compare the results.
4. Other PCI functions replicate DMA behavior, e.g., `pci_unmap_single`. For example, commit `b0eb57cb` contains:

```
- pci_unmap_single(pdev, tbi->dma_addr, tbi->len,  
+ dma_unmap_single(&pdev->dev, tbi->dma_addr, tbi->len,  
                  PCI_DMA_TODEVICE);
```

Extend your semantic patch to implement this transformation. Try to minimize the number of rules.

## Getter and setter functions

Some functions from `include/linux/ide.h`:

```
static inline void *  
ide_get_hwifdata (ide_hwif_t * hwif)  
{  
    return hwif->hwif_data;  
}
```

```
static inline void  
ide_set_hwifdata (ide_hwif_t * hwif, void *data)  
{  
    hwif->hwif_data = data;  
}
```

**Goal:** Replace uses of `hwif->hwif_data` by calls to these function.

## Getter and setter functions: First attempt

@@

```
expression hwif;
```

@@

```
- hwif->hwif_data
```

```
+ ide_get_hwifdata(hwif)
```

@@

```
expression hwif, data;
```

@@

```
- hwif->hwif_data = data
```

```
+ ide_set_hwifdata(hwif, data)
```

# Problems

```
@@ expression hwif; @@  
- hwif->hwif_data  
+ ide_get_hwifdata(hwif)
```

- The rule applies to

```
unsigned long base = (unsigned long)hwif->hwif_data;
```

but also to

```
hwif->hwif_data = NULL;
```

- `ide_get_hwifdata` has prototype:

```
static inline void *ide_get_hwifdata (ide_hwif_t * hwif);
```

The rule transforms **all** `hwif_data` field references.

## Second attempt: Rule order

@@

expression hwif, data;

@@

- hwif->hwif\_data = data

+ ide\_set\_hwifdata(hwif, data)

@@

expression hwif;

@@

- hwif->hwif\_data

+ ide\_get\_hwifdata(hwif)

Applies to 9 code sites, in 2 files.

## Third attempt: Metavariable type constraints

```
@@ ide_hwif_t *hwif; expression data; @@  
- hwif->hwif_data = data  
+ ide_set_hwifdata(hwif, data)
```

```
@@ ide_hwif_t *hwif; @@  
- hwif->hwif_data  
+ ide_get_hwifdata(hwif)
```

Can optionally add `typedef ide_hwif_t;` in the first rule.

- Typedef is needed when the type appears only in a cast.
- Typedef appears only once, where earliest needed.



## Exercise 4

1. Implement all three variants of the semantic patch for introducing the `ide_get_hwifdata` and `ide_set_hwifdata` getter and setter functions.
2. Test each variant on the directory `linux-3.2/drivers/ide`, and compare the results.
3. Reimplement each variant using a disjunction.
4. Compare the results of the disjunction variants to the original implementations.

## Exercise 5

In the case of `pci_map_single` and `dma_map_single`, we could also prefer to convert PCI occurrences of `dma_map_single` to calls to `pci_map_single` (i.e., the reverse transformation).

1. Implement a semantic patch to do this transformation.
2. Test your semantic patch on the directory `linux-3.2/drivers/net/ethernet`.
3. Given the definition of `pci_map_single` in `include/asm-generic/pci-dma-compat.h`, why should the file `ethernet/cadence/macb.c` not be transformed?
4. Check that your semantic patch makes no modification in this file.

# Summary

SmPL features seen so far:

- Metavariables for abstracting over arbitrary expressions.
- Metavariables restricted to particular types.
- Disjunctions.
- Multiple rules.
- Rule ordering.